

# Information Retrieval System

---

## CS6370-NLP Project

---

Team - 20

B Aditi (mm19b020)

R Rishaab Karthik (mm19b046)

```
In [1]: from util import *
        from sentencesegmentation import *
        from tokenisation import *
        from stopwordremoval import *
        from inflectionReduction import *
        from informationretrieval import *
        from evaluation import *
```

```
[nltk_data] Downloading package omw-1.4 to /home/rishaab/nltk_data...
[nltk_data]   Package omw-1.4 is already up-to-date!
[nltk_data] Downloading package punkt to /home/rishaab/nltk_data...
[nltk_data]   Package punkt is already up-to-date!
[nltk_data] Downloading package stopwords to
[nltk_data]   /home/rishaab/nltk_data...
[nltk_data]   Package stopwords is already up-to-date!
[nltk_data] Downloading package wordnet to /home/rishaab/nltk_data...
[nltk_data]   Package wordnet is already up-to-date!
```

```
In [2]: pq=[]
        ss=SentenceSegmentation()
        tk=Tokenization()
        sr=StopwordRemoval()
        st=InflectionReduction()
        ir=InformationRetrieval()
        eva=Evaluation()
```

## Read Queries

```
In [3]: queries_json = json.load(open('./cranfield/' + "cran_queries.json", 'r'))[:]
        query_ids, queries = [item["query number"] for item in queries_json], [item["query"] f
```

## Process queries

```
In [4]: proc_queries=[]
        for q in queries:
            ssq=ss.punkt(q)
            tkq=tk.pennTreeBank(ssq)
            srq=sr.fromList(tkq)
            stq=st.reduce(srq)
            proc_queries.append(stq)
```

## Read documents

```
In [5]: docs_json = json.load(open('./cranfield/' + "cran_docs.json", 'r'))[:]
```

```
doc_ids, docs = [item["id"] for item in docs_json], [item["body"] for item in docs_json]
```

## Process Documents

```
In [6]: proc_docs=[]
        for d in docs:
            ssd=ss.punkt(d)
            tkd=tk.pennTreeBank(ssd)
            srd=sr.fromList(tkd)
            std=st.reduce(srd)
            proc_docs.append(std)
```

## TF\_IDF Matrix

```
In [7]: tfidf_docs_df=ir.buildIndex(proc_docs,doc_ids)
```

```
In [11]: tfidf_docs_df.shape
```

```
Out[11]: (1400, 6611)
```

```
In [8]: tfidf_queries_df=ir.tfidfq(proc_queries)
```

## Latent Semantic analysis

```
In [28]: lsadocs,lsaw2v=ir.lsa_doc(tfidf_docs_df,500)
```

```
In [29]: lsaqueries,lsa2v=ir.lsa_query(tfidf_queries_df)
```

## Word2Vec Embedding on LSA

```
In [30]: doc_lsa_w2v=ir.w2v(lsaw2v)
```

```
In [31]: q_lsa_w2v=ir.w2v(lsa2v)
```

## Word2Vec embedding standalone

```
In [32]: doc_w2v=ir.w2v(tfidf_docs_df)
        q_w2v=ir.w2v(tfidf_queries_df)
```

## Different Models Evaluated

```
In [35]: doc_ord_mod1=ir.rank(doc_ids,tfidf_docs_df,tfidf_queries_df)
```

```
In [34]: doc_ord_mod2=ir.rank(doc_ids,lsadocs,lsaqueries)
```

```
In [33]: doc_ord_mod3=ir.rank(doc_ids,doc_lsa_w2v,q_lsa_w2v)
```

```
In [36]: doc_ord_mod4=ir.rank(doc_ids,doc_w2v,q_w2v)
```

```
In [37]: ev=Evaluation()
```

```
In [38]: grels = json.load(open('./cranfield/' + "cran_grels.json", 'r'))[:]
```

## Plain VSM TF\_IDF Model

```
In [42]: ev.meanAveragePrecision(doc_ord_mod1,query_ids,qrels,10)
```

```
Out[42]: 0.32691764858923683
```

## LSA Model

```
In [41]: ev.meanAveragePrecision(doc_ord_mod2,query_ids,qrels,10)
```

```
Out[41]: 0.3328471536867305
```

## LSA + Word2Vec

```
In [40]: ev.meanAveragePrecision(doc_ord_mod3,query_ids,qrels,10)
```

```
Out[40]: 0.2512909672215898
```

## Word2Vec

```
In [44]: ev.meanAveragePrecision(doc_ord_mod4,query_ids,qrels,10)
```

```
/home/rishaab/Desktop/SEM8/CS6370/project/Code/evaluation.py:32: FutureWarning: The behavior of `series[i:j]` with an integer-dtype index is deprecated. In a future version, this will be treated as *label-based* indexing, consistent with e.g. `series[i]` lookups. To retain the old behavior, use `series.iloc[i:j]`. To get the future behavior, use `series.loc[i:j]`.
  for docID in query_doc_IDs_ordered[:k]:
```

```
Out[44]: 0.20479112536927754
```

## Finding optimal number of components for LSA and LSA+Word2Vec based on MAP

```
In [45]: def lsa_model(k):
  lsadocs=ir.lsa_doc(tfidf_docs_df,k)[0]
  lsaqueries=ir.lsa_query(tfidf_queries_df)[0]
  doc_ord=ir.rank(doc_ids,lsadocs,lsaqueries)
  return doc_ord,ev.meanAveragePrecision(doc_ord,query_ids,qrels,1)
```

```
In [46]: def lsa_model2(k):
  lsadocs=ir.lsa_doc(tfidf_docs_df,k)[1]
  lsaqueries=ir.lsa_query(tfidf_queries_df)[1]
  lsa_w2v=ir.w2v(lsadocs)
  q_w2v=ir.w2v(lsaqueries)
  doc_ord=ir.rank(doc_ids,lsa_w2v,q_w2v)
  return doc_ord,ev.meanAveragePrecision(doc_ord,query_ids,qrels,1)
```


```
In [47]: order=[]
  MAP=[]
  for i in range(100,3100,100):
    o,m = lsa_model(i)
    order.append(o)
    MAP.append(m)
```



[illegible]



Number of Latent Variables Vs MAP for LSA



```
In [49]: order2=[]  
MAP2=[]  
for i in range(100,3100,100):  
    o,m = lsa_model2(i)  
    order2.append(o)  
    MAP2.append(m)
```








[illegible]



Number of Latent Variables Vs MAP for LSA+W2V



## Functions to plot different Evaluation metrics

```
In [51]: def eval_met(doc_ord_mod1):
    prec=[]
    meanap=[]
    meanFscore=[]
    meanrecall=[]
    nDCG=[]
    for i in range(1,20):
        prec.append(ev.meanPrecision(doc_ord_mod1,query_ids,qrels,i))
        meanap.append(ev.meanAveragePrecision(doc_ord_mod1,query_ids,qrels,i))
        meanFscore.append(ev.meanFscore(doc_ord_mod1,query_ids,qrels,i))
        meanrecall.append(ev.meanRecall(doc_ord_mod1,query_ids,qrels,i))
        nDCG.append(ev.meanNDCG(doc_ord_mod1,query_ids,qrels,i))
    modell=pd.DataFrame({'Precision':prec,'MaP':meanap,"meanFscore":meanFscore,'meanr
    return modell
```

```
In [59]: def plot_model(df,title):
    plt.figure(figsize=[12,10])
    for i in df.columns:
        plt.plot(range(1,20),df[i],label=i)
        plt.legend()
        plt.xlabel("k")
        plt.title(title)
```

## VSM-TF\_IDF

```
In [56]: modell=eval_met(doc_ord_mod1)
```


```
/home/rishaab/Desktop/SEM8/CS6370/project/Code/evaluation.py:32: FutureWarning: The behavior of `series[i:j]` with an integer-dtype index is deprecated. In a future version, this will be treated as *label-based* indexing, consistent with e.g. `series[i]` lookups. To retain the old behavior, use `series.iloc[i:j]`. To get the future behavior, use `series.loc[i:j]`.
```

```
for docID in query_doc_IDs_ordered[:k]:
```

```
/home/rishaab/Desktop/SEM8/CS6370/project/Code/evaluation.py:118: FutureWarning: The behavior of `series[i:j]` with an integer-dtype index is deprecated. In a future version, this will be treated as *label-based* indexing, consistent with e.g. `series[i]` lookups. To retain the old behavior, use `series.iloc[i:j]`. To get the future behavior, use `series.loc[i:j]`.
```

```
for docID in query_doc_IDs_ordered[:k]:
```

```
In [60]: plot_model(model1, 'Evaluation Metrics VSM-TFIDF')
```




## LSA

```
In [61]: doc_ord_opt_lsa=order[2]
eval_model_opt_lsa=eval_met(doc_ord_opt_lsa)
```

```
In [62]: plot_model(eval_model_opt_lsa, 'Evaluation Metrics Optimised LSA model')
```

Evaluation Metrics Optimised LSA model




## LSA + Word2Vec

```
In [65]: eval_lsa_w2v_opt=eval_met(order2[2])
```

```
In [66]: plot_model(eval_lsa_w2v_opt,'Evaluation Metrics Optimised LSA+W2V model')
```

Evaluation Metrics Optimised LSA+W2V model



## Word 2 Vec


```
In [69]: w2v_eval=eval_met(doc_ord_mod4)
plot_model(w2v_eval,'Evaluation Metrics W2V Model')
```

/home/rishaab/Desktop/SEM8/CS6370/project/Code/evaluation.py:32: FutureWarning: The behavior of `series[i:j]` with an integer-dtype index is deprecated. In a future version, this will be treated as \*label-based\* indexing, consistent with e.g. `series[i]` lookups. To retain the old behavior, use `series.iloc[i:j]`. To get the future behavior, use `series.loc[i:j]`.

```
for docID in query_doc_IDs_ordered[:k]:
```

/home/rishaab/Desktop/SEM8/CS6370/project/Code/evaluation.py:118: FutureWarning: The behavior of `series[i:j]` with an integer-dtype index is deprecated. In a future version, this will be treated as \*label-based\* indexing, consistent with e.g. `series[i]` lookups. To retain the old behavior, use `series.iloc[i:j]`. To get the future behavior, use `series.loc[i:j]`.

```
for docID in query_doc_IDs_ordered[:k]:
```



## Results

```
In [82]: results=pd.concat([model1.iloc[9:10],eval_model_opt_lsa[9:10],eval_lsa_w2v_opt[9:10],
results['Model']=['TFIDF',"LSA","LSA+W2V","W2V"]
results.set_index('Model',inplace=True)
```

```
In [84]: print('The Evaluation Metric values @ k=10')
print(results)
```

The Evaluation Metric values @ k=10

	Precision	MaP	meanFscore	meanrecall	MeanNDCG
Model					
TFIDF	0.295556	0.326918	0.321014	0.420230	0.003142
LSA	0.304444	0.338259	0.331191	0.434138	0.003142
LSA+W2V	0.237778	0.261831	0.263169	0.349854	0.003142
W2V	0.200000	0.204791	0.221611	0.301034	0.003142